

Lecture 11: Analyzing data II – Grouping and analyzing

Topic – Summary functions

Reread in the csv file we covered last lecture.

```
>>> import pandas as pd
>>> dt = pd.read_csv('021821_medium_threshold2.csv', index_col = 0)
```

Dataframes come with summary statistics functions, which summarize the value for a row or column. For example, if you wanted to calculate the average gene expression in each cell type, you could do the following (use drop to get rid of the non-numerical columns).

```
>>> dt.drop(columns = ['gene_name', 'Wormbase_ID']).mean()
ADA      70.876487
ADE      71.643830
ADF      72.214044
ADL      71.873994
AFD      71.160180
...
VB01     72.527223
VB02     72.368139
VC       71.968446
VC_4_5   72.100255
VD_DD    70.920625
Length: 128, dtype: float64
```

Other summary functions include median(), min(), max(), count(), std(), kurtosis(). There is also a describe method that includes a number of these functions:

```
>>> dt.drop(columns = ['gene_name', 'Wormbase_ID']).describe()
          ADA      ADE      ADF      ADL      ...      VB02
VC          VC_4_5      VD_DD
count  13669.000000  13669.000000  13669.000000  13669.000000  ...  13669.000000
13669.000000  13669.000000  13669.000000
mean    70.876487    71.643830    72.214044    71.873994  ...    72.368139
71.968446    72.100255    70.920625
std     835.806300   1310.171641   2403.896102   2922.675367  ...    823.131334
1357.371236   1464.616460    964.106266
min      0.000000      0.000000      0.000000      0.000000  ...      0.000000
0.000000      0.000000      0.000000
25%      0.000000      0.000000      0.000000      0.000000  ...      0.000000
0.000000      0.000000      0.000000
50%      0.000000      0.000000      1.626997      1.072438  ...      1.102683
0.000000      0.000000      0.000000
75%     14.118176    14.698212    16.613110    16.421330  ...    24.470680
17.023660    14.461254    24.199716
max     46308.923938   98000.306128   202697.711797   336350.854760  ...   56469.718483
92023.392079   89161.640000   80754.488434
```

You can also run these on rows, using the axis argument.

```
dt.drop(columns = ['gene_name', 'Wormbase_ID']).mean(axis = 1)
1      8926.858079
2     1060.437786
3      856.869660
4     1519.028523
5     4458.778125
...
13665      0.674315
13666      4.338739
```

```
13667      1.186117
13668      2.618098
13669      1.182652
Length: 13669, dtype: float64
```

Topic – Creating and modifying a dataframe

It is also possible to create a dataframe from other objects. Typically, this easiest using a dictionary, where the key is the name of the column and the value is a list of data. For example:

```
pd.DataFrame({'Test': [1,2,3,4], 'OtherData': [4,3,2,1]})
```

	Test	OtherData
0	1	4
1	2	3
2	3	2
3	4	1

Notice that an index is automatically created for you.

If you want to specify the index, you can use the index argument and supply it with a list that is the same size as the other data like so:

```
>>> pd.DataFrame({'Test': [1,2,3,4], 'OtherData': [4,3,2,1]}, index = ['A','B','C','D'])
```

	Test	OtherData
A	1	4
B	2	3
C	3	2
D	4	1

If you would like to modify the dataframe, you can modify a single element pretty straightforwardly. Recall how to use loc to access an individual cell and then set it equal to whatever you want:

```
>>> dt = pd.DataFrame({'Test': [1,2,3,4], 'OtherData': [4,3,2,1]}, index = ['A','B','C','D'])
>>> dt.loc['A','Test'] = 3
>>> dt
```

	Test	OtherData
A	3	4
B	2	3
C	3	2
D	4	1

You can also add an entire column by accessing the new columns name and setting it equal to a list:

```
>>> dt['NewData'] = [3,4,2,5]
>>> dt
```

	Test	OtherData	NewData
A	3	4	3
B	2	3	4
C	3	2	2
D	4	1	5

Or if you want to create a new column based upon existing data you can use apply to apply a function to every element of a column.

```
>>> import math
>>> dt['NewDataSquareRoot'] = dt['NewData'].apply(math.sqrt)
>>> dt
```

	Test	OtherData	NewData	NewDataSquareRoot
A	3	4	3	1.732051

B	2	3	4	2.000000
C	3	2	2	1.414214
D	4	1	5	2.236068

Or even get more complicated:

```
dt['CombinedData'] = dt['NewData'].apply(math.sqrt) + dt.OtherData
>>> dt
```

	Test	OtherData	NewData	NewDataSquareRoot	CombinedData
A	3	4	3	1.732051	5.732051
B	2	3	4	2.000000	5.000000
C	3	2	2	1.414214	3.414214
D	4	1	5	2.236068	3.236068

Topic – Grouping data

Data in a dataframe is organized in a tabular format, which is often most easily thought of as a collection of records. Each ROW consists of data for one record. The type of data that is stored is specified by the columns. For example, you might have a large data file containing information for every flight that flew in the US. Each record is UNIQUE. In other words, no two rows should be exactly the same as each flight is unique.

However, each column is not necessarily unique. For example, for the flight example above, you might store the airport that the flight left from. Since many flights leave from a given airport, there are many rows that share the same origin airport information. An example is shown below.

Now, sometimes you want to GROUP records together that share the same values. For example, let's say I wanted to find the airport that had the longest delays. To do that I need to average all of the delays together for EVERY flight that leaves from that airport.

Step 1: identify all the rows that have airport 1 together and group them together. In your mind you can think of it as putting them all in the same bucket. The number of buckets is determined by the number of unique values in a column (i.e. the number of airports in the US). Each bucket also contains multiple records (i.e. the number of flights that left from a particular airport).

Your next goal is to take all of the records in a bucket and condense them down to a single number. For example, you might want to find the average time delay for each airport. In this case you would take the AVERAGE for all the records in the bucket. Or you might want to find the maximum delay. Then you would take the MAX of the records.

origin	depdelay
airport 1	10
airport 1	15
airport 2	100
airport 3	320
airport 1	20
airport 2	120
airport 2	140
airport 1	15

airport 1 → avg: 15

airport 2 → avg: 120

airport 3 → avg: 320

Your next goal is to take all of the records in a bucket and condense them down to a single number. For example, you might want to find the average time delay for each airport. In this case you would take the AVERAGE for all the records in the bucket. Or you might want to find the maximum delay. Then you would take the MAX of the records. Let's see an example of this using the WorldDemographics.csv file:

```
>>> import pandas as pd
>>> dt = pd.read_csv('WorldDemographics.csv', sep = ',', index_col = 0)
>>> dt
```

	PopulationID	country_code	Level	continent_code	Age	#Alive
0	Afghanistan	4	Country	5501	0	1134501
1	Afghanistan	4	Country	5501	1	1134501
2	Afghanistan	4	Country	5501	2	1134501
3	Afghanistan	4	Country	5501	3	1124250
4	Afghanistan	4	Country	5501	4	1113998
...
20296	Zimbabwe	716	Country	910	96	163
20297	Zimbabwe	716	Country	910	97	53
20298	Zimbabwe	716	Country	910	98	41
20299	Zimbabwe	716	Country	910	99	30
20300	Zimbabwe	716	Country	910	100	19

[20301 rows x 6 columns]

Notice that there are many records or rows that share the sample Population ID. What if we wanted to find the total population of a country? We could do that using a groupby followed by a summary command.

```
>>> country_data = dt.groupby(['PopulationID'])
>>> type(country_data)
<class 'pandas.core.groupby.generic.DataFrameGroupBy'>
```

We use the groupby method of the DataFrame class, passing it the name of the column we want to group by (we could also send it a list of column names if we wanted to group on two columns).

Now we need to specify what function to use to combine all the records:

```
>>> country_data.sum()
```

	country_code	continent_code	Age	#Alive
PopulationID				
Afghanistan	404	555601	5050	38928204
Albania	808	93425	5050	2877640
Algeria	1212	92112	5050	43850203
Angola	2424	92011	5050	32866164
Antigua and Barbuda	2828	92415	5050	97878
...
Viet Nam	71104	92920	5050	97342743
Western Sahara	73932	92112	5050	597282
Yemen	89587	93122	5050	29825815
Zambia	90294	91910	5050	18383891
Zimbabwe	72316	91910	5050	14862850

[201 rows x 4 columns]

```
>>> country_data.mean()
```

	country_code	continent_code	Age	#Alive
PopulationID				
Afghanistan	4.0	5501.0	50.0	385427.762376
Albania	8.0	925.0	50.0	28491.485149
Algeria	12.0	912.0	50.0	434160.425743
Angola	24.0	911.0	50.0	325407.564356
Antigua and Barbuda	28.0	915.0	50.0	969.089109

```

...
Viet Nam          704.0          920.0  50.0  963789.534653
Western Sahara    732.0          912.0  50.0    5913.683168
Yemen             887.0          922.0  50.0  295305.099010
Zambia            894.0          910.0  50.0  182018.722772
Zimbabwe          716.0          910.0  50.0  147156.930693

```

[201 rows x 4 columns]

```

>>> country_data.min()

```

	country_code	Level	continent_code	Age	#Alive
PopulationID					
Afghanistan	4	Country	5501	0	39
Albania	8	Country	925	0	51
Algeria	12	Country	912	0	963
Angola	24	Country	911	0	38
Antigua and Barbuda	28	Country	915	0	4
...
Viet Nam	704	Country	920	0	19645
Western Sahara	732	Country	912	0	1
Yemen	887	Country	922	0	78
Zambia	894	Country	910	0	5
Zimbabwe	716	Country	910	0	19

Notice this works on all of the data columns that have data that can handle the function (i.e. you can't sum a string but you can take the minimum of it). If you want to only act on specific columns, use the agg method and specify the column/function combo using a dictionary:

```

>>> country_data.agg({'Age':min, '#Alive':sum})

```

	Age	#Alive
PopulationID		
Afghanistan	0	38928204
Albania	0	2877640
Algeria	0	43850203
Angola	0	32866164
Antigua and Barbuda	0	97878
...
Viet Nam	0	97342743
Western Sahara	0	597282
Yemen	0	29825815
Zambia	0	18383891
Zimbabwe	0	14862850

[201 rows x 2 columns]

The numpy module contains a large number of functions that are useful for analyzing arrays of data. We will cover those later. For now just remember that you will probably want to use numpy functions in the future for more complex summary operations.

```

>>> country_dt = country_data.agg({'Age':min, '#Alive':sum})

```

```

>>> country_dt

```

	Age	#Alive
PopulationID		
Afghanistan	0	38928204
Albania	0	2877640
Algeria	0	43850203
Angola	0	32866164
Antigua and Barbuda	0	97878
...
Viet Nam	0	97342743
Western Sahara	0	597282
Yemen	0	29825815

Zambia	0	18383891
Zimbabwe	0	14862850

[201 rows x 2 columns]

```
>>> country_dt.loc['Angola','#Alive']  
32866164
```